



# Systems and Internet Infrastructure Security

Network and Security Research Center  
Department of Computer Science and Engineering  
Pennsylvania State University, University Park PA

## CSE543 - Introduction to Computer and Network Security Module: Operating System Security

Professor Trent Jaeger

- So, you have built an operating system that enables user-space processes to access hardware resources
  - Thru various abstractions: files, pages, devices, etc.
- Now, you want your operating system to enforce **security requirements** for your application processes
  - What do you do?



- We learned about a few things that will help you
- Your OS must implement a
  - ▶ (Mandatory) Protection system
- That can enforce a
  - ▶ MAC policy
- How do we implement such an OS mechanism?
  - ▶ Multics
  - ▶ Linux Security Modules



# Access Policy Enforcement

- A protection system uses a *reference validation mechanism* to produce and evaluate authorization queries
  - ▶ **Interface**: Mediate *security-sensitive operations* by building authorization queries to evaluate
  - ▶ **Module**: Determine relevant *protection state* entry (ACLs, capabilities) to evaluate authorization query
  - ▶ **Manage**: Install protection state entries and reason about *labeling and transition states*
- How do we know whether a reference validation mechanism is correct?

# Security-Sensitive Operations

- Broadly, operations that **enable interaction among processes** that violate secrecy, integrity, availability
- Which of these are security-sensitive? Why?
  - ▶ Read a file (*read*)
  - ▶ Get the process id of a process (*getpid*)
  - ▶ Read file metadata (*stat*)
  - ▶ Fork a child process (*fork*)
  - ▶ Get the metadata of a file you have already opened? (*fstat*)
  - ▶ Modify the data segment size? (*brk*)
- Require protection for all of **CIA?**

# Reference Monitor

- Defines a set of requirements on **reference validation mechanisms**
  - To enforce access control policies correctly
- **Complete mediation**
  - The *reference validation mechanism* must always be invoked (before executing security-sensitive operations)
- **Tamperproof**
  - The *reference validation mechanism* must be tamperproof
- **Verifiable**
  - The *reference validation mechanism* must be small enough to be subject to analysis and tests, the completeness of which can be assured

# Multiprocessor Systems

- Major Effort: *Multics*
  - ▶ Multiprocessing system -- developed many OS concepts
    - Including security
  - ▶ Begun in 1965
    - Research continued into the mid-70s
  - ▶ Used until 2000
  - ▶ Initial partners: MIT, Bell Labs, GE (replaced by Honeywell)
  - ▶ *Other innovations*: hierarchical filesystems, dynamic linking
- Subsequent proprietary system, *SCOMP*, became the basis for secure operating systems design (XTS-400)



# Multics Goals


- Secrecy
  - Multilevel security
- Integrity
  - Rings of protection
- Resulting system is considered a high point in secure systems design





# Protection Rings


- Successively less-privileged “domains”
- Modern CPUs support 4 rings
  - ▶ Use 2 mainly: Kernel and user
- Intel x86 rings
  - ▶ Ring 0 has kernel
  - ▶ Ring 3 has application code



- Example: Multics (64 rings in theory, 8 in practice)


# What Are Protection Rings?

- Coarse-grained, Hardware Protection Mechanism
- Boundary between Levels of Authority
  - Most privileged -- ring 0
  - Monotonically less privileged above
- Fundamental Purpose
  - **Protect system integrity**
    - Protect kernel from services
    - Protect services from apps
    - So on...




# Protection Ring Rules

- Program cannot call code of *higher privilege* directly
  - ▶ Gate is a special memory address where lower-privilege code can call higher
    - Enables OS to control where applications call it (system calls)



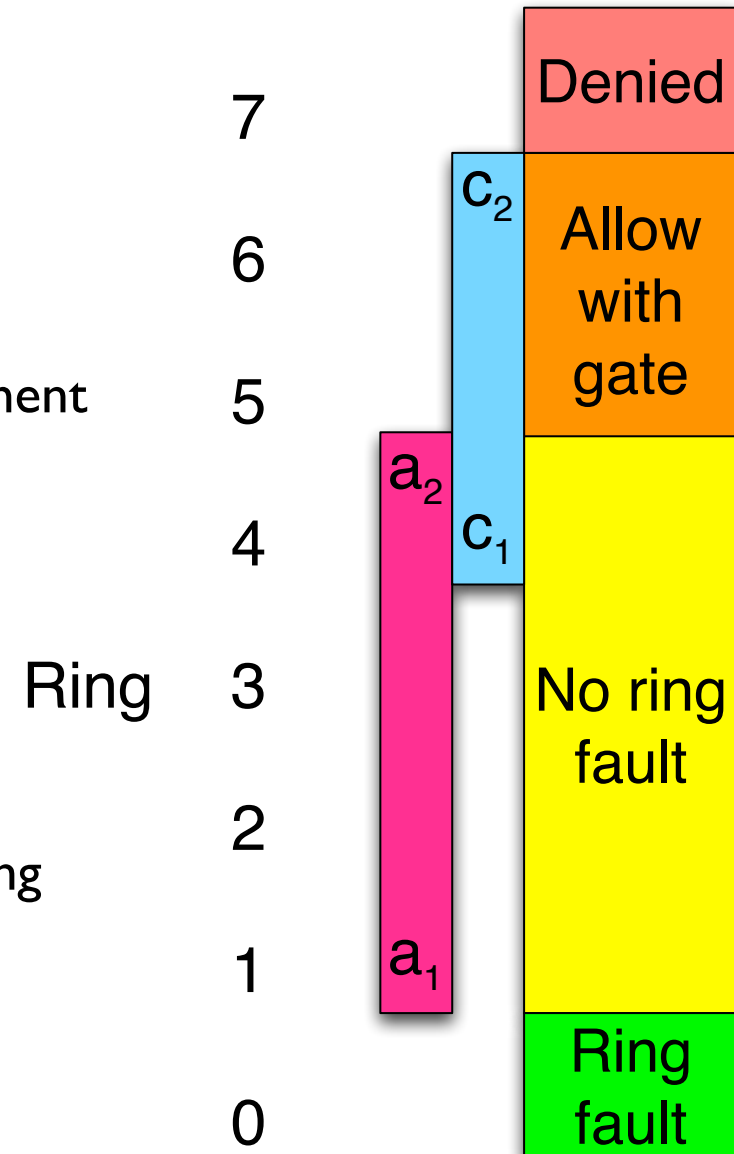
# Multics Interpretation

- Kernel resides in ring 0
- Process runs in a ring  $r$ 
  - Access based on current ring
- Process accesses data (segment)
  - Each data segment has an *access bracket*:  $(a_1, a_2)$ 
    - $a_1 \leq a_2$
  - Describes read and write access to segment
    - $r$  is the current ring
    - $r \leq a_1$ : access permitted
    - $a_1 < r \leq a_2$ :  $r$  and  $x$  permitted;  $w$  denied
    - $a_2 < r$ : all access denied



# Multics Interpretation (con't)


- Also different procedure segments
  - with *call brackets*:  $(c1, c2)$ ,  $c1 \leq c2$
  - and access brackets  $(a1, a2)$
  - The following must be true ( $a2 == c1$ )
  - Rights to execute code in a new procedure segment
    - $r < a1$ : access permitted with ring-crossing fault
    - $a1 \leq r \leq a2 = c1$ : access permitted and no fault
    - $a2 < r \leq c2$ : access permitted through a valid gate
    - $c2 < r$ : access denied
- What's it mean?
  - case 1: ring-crossing fault changes procedure's ring
    - increases from  $r$  to  $a1$
  - case 2: keep same ring number
  - case 3: gate checks args, decreases ring number
- Target code segment defines the new ring



# Examples

- Process in ring 3 accesses data segment
  - access bracket: (2, 4)
  - What operations can be performed?
- Process in ring 5 accesses same data segment
  - What operations can be performed?
- Process in ring 5 accesses procedure segment
  - access bracket (2, 4)
  - call bracket (4, 6)
  - Can call be made?
  - How do we determine the new ring?
  - Can new procedure segment access the data segment above?

# Now forward to UNIX ...



# UNIX Security Limitations

- *Circa 2000 Problems*
  - ▶ Discretionary access control
  - ▶ Setuid root processes
  - ▶ Network-facing daemons vulnerable
- What can we do?



# UNIX Security Limitations

- *Circa 2000 Problems*
  - ▶ Discretionary access control
  - ▶ Setuid root processes
  - ▶ Network-facing daemons vulnerable
- What can we do?
  - ▶ Reference validation mechanism that satisfies reference monitor concept
  - ▶ Protection system with mandatory access control (mandatory protection system)

# Linux Security Modules

- Reference validation mechanism for Linux
  - Upstreamed in Linux 2.6
  - Support modular enforcement - you choose
    - SELinux, AppArmor, POSIX Capabilities, SMACK, ...
- 150+ **authorization hooks**
  - Mediate security-sensitive operations on
    - Files, dirs/links, IPC, network, semaphores, shared memory, ...
  - Variety of operations per data type
    - Control access to read of file data and file metadata separately
- Hooks are **restrictive**


# LSM & Reference Monitor

- Does LSM satisfy reference monitor concept?

# LSM & Reference Monitor

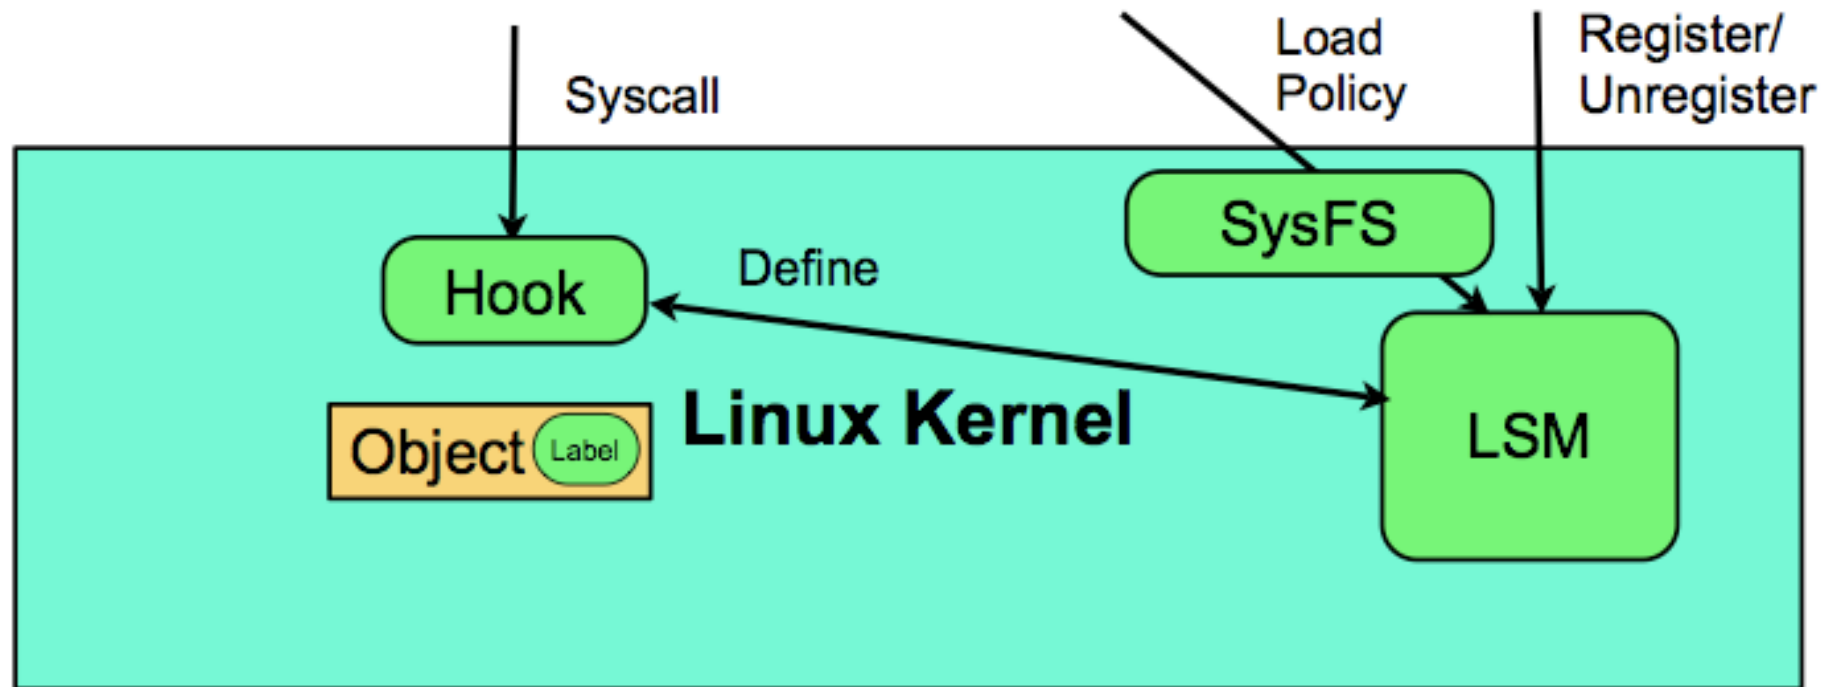
- Does LSM satisfy reference monitor concept?
  - ▶ **Tamperproof**
    - Can **MAC policy** be tampered?
    - Can **kernel** be tampered?

# Linux Security Modules



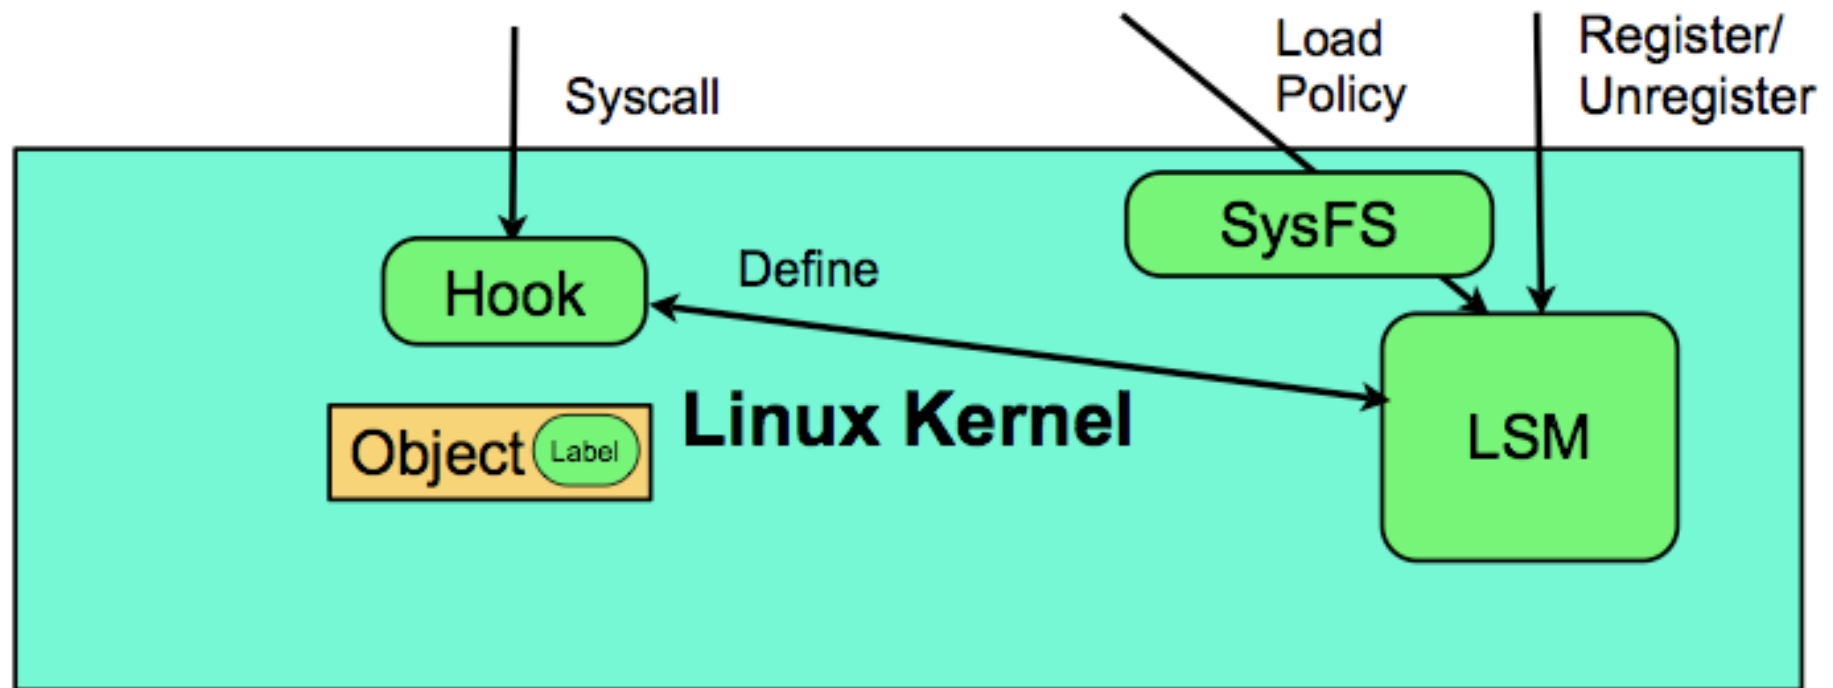
- Register (install) module
- Load policy (open and write to special file)
- Produce authorization queries at hooks

# Linux Security Modules



- Attacks on “register”
- Attacks on “install policy”
- Attacks on “system calls”

# Linux Security Modules



- To prevent attacks on registration
- And attacks on function pointers of LSM
- **LSMs are now statically compiled into the kernel**

# LSM & Reference Monitor

- Does LSM satisfy reference monitor concept?
  - ▶ Tamperproof
    - Can MAC policy be tampered?
    - Can kernel be tampered?
  - ▶ Verifiable
    - How large is kernel?
    - Can we perform **complete** testing?



# LSM & Reference Monitor

- Does LSM satisfy reference monitor concept?
  - ▶ Tamperproof
    - Can MAC policy be tampered?
    - Can kernel be tampered? By network threats?
  - ▶ Verifiable
    - How large is kernel?
    - Can we perform complete testing?
  - ▶ Complete Mediation
    - What is a **security-sensitive operation**?
    - Do we mediate all paths to such operations?

# Linux Security Modules

Security check function

```
linux/fs/read_write.c:  
  
ssize_t vfs_read(...) {  
    ...  
    ret = security_file_permission(file, ...);  
    if (!ret) { ...  
        ret = file->f_op->read(file, ...); ...  
    }  
    ...  
}
```

Security sensitive operation

# LSM & Complete Mediation

- What is a security-sensitive operation?
  - ▶ **Instructions?** Which?
  - ▶ **Structure member accesses?** To what data?
  - ▶ **Data types** whose instances may be controlled
    - Inodes, files, IPCs, tasks, ...
- Approaches
  - ▶ **Mediation:** Check that authorization hook dominates all control-flow paths to structure member access on security-sensitive data type
  - ▶ **Consistency:** Check that every structure member access that is mediated once is always mediated
    - Several bugs found - some years later

# LSM & Complete Mediation

- Static analysis of Zhang, Edwards, and Jaeger [USENIX Security 2002]
  - ▶ Based on a tool called CQUAL
- Found a **TOCTTOU** bug
  - ▶ Authorize *filp* in *sys\_fcntl*
  - ▶ But pass *fd* again to *fcntl\_getlk*
- Many supplementary analyses were necessary to support CQUAL

```
/* from fs/fcntl.c */
long sys_fcntl(unsigned int fd,
               unsigned int cmd,
               unsigned long arg)
{
    struct file * filp;
    ...
    filp = fget(fd);
    ...

    err = security_ops->file_ops
        ->fcntl(filp, cmd, arg);
    ...
    err = do_fcntl(fd, cmd, arg, filp);
    ...
}

static long
do_fcntl(unsigned int fd,
         unsigned int cmd,
         unsigned long arg,
         struct file * filp) {
    ...
    switch(cmd){
        ...
        case F_SETLK:
            err = fcntl_setlk(fd, ...);
            ...
        }
    ...
}

/* from fs/locks.c */
fcntl_getlk(fd, ...) {
    struct file * filp;
    ...

    filp = fget(fd);

    /* operate on filp */
    ...
}
```

Figure 8: Code path from Linux 2.4.9 containing an exploitable type error.

# LSM Enforcement

- Several LSMs have been deployed
  - ▶ Most prominent: AppArmor, SELinux, Smack, TOMOYO
- The most comprehensive is [SELinux](#)
  - ▶ Used by RedHat Fedora and some others

# LSM Enforcement

- Several LSMs have been deployed
  - ▶ Most prominent: AppArmor, SELinux, Smack, TOMOYO
- The most comprehensive is SELinux
  - ▶ Created by the NSA - Result of many years work
  - ▶ Used by RedHat Fedora and some others



# SELinux Policy Rules

- SELinux Rules express an MPS
  - ▶ *Protection state* – ALLOW subject-label object-label ops
  - ▶ *Labeling state* – Assign new objects labels on creation
  - ▶ *Transition state* – Define how a process may change label
- All are defined explicitly
  - ▶ Tens of thousands of rules are necessary for a standard Linux distribution
    - Remember, we are ignoring user processes too (other than confining them relative to the system)
- Enforces a Least Privilege Policy

# SELinux Transition State

- For user to run passwd program
  - Only passwd should have permission to modify */etc/shadow*
- Need permission to execute the passwd program
  - *allow user\_t passwd\_exec\_t:file execute* (user can exec */usr/bin/passwd*)
  - *allow user\_t passwd\_t:process transition* (user gets passwd perms)
- Must transition to passwd\_t from user\_t
  - *allow passwd\_t passwd\_exec\_t:file entrypoint* (run w/ passwd perms)
  - *type\_transition user\_t passwd\_exec\_t:process passwd\_t*
- Passwd can the perform the operation
  - *allow passwd\_t shadow\_t:file {read write}* (can edit passwd file)



# Take Away

- **Goal:** Build authorization into operating systems
  - ▶ Multics and Linux
- **Requirements:** Reference monitor
  - ▶ Satisfy reference monitor concept
- Multics
  - ▶ Hierarchical Rings for Protection
  - ▶ Call/Access Bracket Policies (in addition to MLS)
- Linux
  - ▶ Did not enforce **security** (*DAC, Setuid, root daemons*)
  - ▶ So, the Linux Security Modules framework was added
  - ▶ Approximates reference monitor assuming network threats only -- some challenges in ensuring complete mediation